

Building and Evaluating a Retrieval-Augmented Generation System on FinanceBench

Building and Evaluating a Retrieval-Augmented Generation System on FinanceBench

Karthik Reddy Changal, Chaithanya Anugu, Nithin Sujith Nair

Abstract

We build and evaluate a retrieval-augmented generation (RAG) pipeline that answers questions about financial filings using only a closed collection of SEC documents, on free models that run locally. We compare configurations across chunk strategy, retrieved-passage count, embedding model, generation model, retrieval filtering, cross-encoder reranking, and prompting. On the 114 questions whose source document is available, the best configuration restricts retrieval to the company named in the question and reorders a deep candidate pool with a cross-encoder, reaching Recall@5 0.52 against 0.46 for pure dense search. The headline result is methodological: measured by Recall@5 the company filter supplies 88% of the gain and reranking only 12%, but measured by whether the system answers correctly the shares are almost exactly reversed, because reranking raises how often the right filing reaches the generator without moving the token-overlap score the retrieval metric is built on. Reranking doubles the number of correct answers; a BM25 hybrid, a three-times-larger generator and a prompt written to reduce refusals all fail to help. Generation remains the weaker stage, and the system declines to answer rather than inventing figures on roughly two-thirds of questions. We report every metric over both the full 150-question set and the answerable subset, because a fifth of the FinanceBench corpus no longer downloads, and we treat that coverage gap as a result rather than hiding it.

1. Introduction

1.1 Problem statement

Financial filings are long and table-heavy; a 10-K runs past 100 pages, and the target fact (a capital expenditure, a segment’s growth) usually sits in one row of one table. Answering a question has two parts: finding the passage with the fact, and reading it correctly. General model knowledge does not help, because the answer is specific to one company and year and must come from the document. Retrieval-augmented generation (RAG) matches this: a retriever selects a few relevant passages and a language model writes a cited answer from them.

1.2 Related work

RAG pairs a dense retriever with a generator so a model answers from an external corpus rather than its parameters (Lewis et al., 2020). Passages and query are embedded into one vector space and compared by similarity. We use FinanceBench (Islam et al., 2023), which gives each question a gold answer and the exact evidence span it draws on, so we can measure both answer quality and whether retrieval found the right passage.

Retrieval is scored by whether the gold evidence appears among the retrieved passages (Recall@k, MRR); generation against the reference answer (ROUGE, semantic similarity, or an LLM judge such as RAGAS). We do not use an LLM judge: it adds a second large model and ties scores to the judge’s quirks. We use lexical and semantic overlap plus a citation-validity check instead.

1.3 Contributions

We build a RAG question-answering pipeline that runs end to end on free, local models, and ablate seven design decisions (chunk strategy, retrieved-passage count, embedding model, generation model, retrieval filtering, cross-encoder reranking, and prompting) one variable at a time. Because much of the FinanceBench corpus no longer downloads, we report every metric over both the full question set and the answerable subset, and treat the gap as a result. Two findings follow. Retrieval, not generation, is the binding constraint: the best configuration finds the gold evidence for barely half the answerable questions. And the retrieval metric itself is an unreliable guide to which changes help, ranking the two interventions we tested in the opposite order to their effect on answer quality, which is a caution for anyone tuning a RAG system against Recall@k alone.

2. Methodology

2.1 Dataset

We use FinanceBench, distributed as two JSONL files linked by doc_name. The document catalogue lists 360 filings; the question file holds 150 questions with gold answers and the evidence text each answer is drawn from.

The filings span several document types (269 10-Ks, 30 8-Ks, 29 earnings releases, 27 10-Qs, 6 annual reports) across eight GICS sectors, led by Information Technology (80), Consumer Discretionary (77), and Consumer Staples (63). The 150 questions divide evenly into three types: metrics-generated, domain-relevant, and novel-generated (50 each). By reasoning type, 43 need numerical reasoning, 31 are information extraction, and the rest mix numerical and logical reasoning. Answers range from single dollar figures (“\$1577.00”) to short explanatory sentences.

Corpus coverage. Only 282 of the 360 catalogued PDFs still download; the other 78 URLs return timeouts, 404s, or 403s from the filers’ investor-relations hosts. A further set of downloaded files are corrupt (for example the AMD 10-Ks, which lack a valid PDF root object) or are near-empty stubs, so text extraction succeeds for 263 documents. After extraction, 114 of the 150 questions reference a document that yields usable text; the remaining 36 point at a missing or corrupt filing and cannot be answered from the corpus. We therefore report each metric twice: over all 150 questions, and over this 114-question answerable subset. The first figure describes the system against the corpus as it exists today; the second isolates the quality of retrieval and generation from missing inputs. Figure 1 traces the decay from the full manifest to the documents that yield usable text.




Figure 1

Figure 1. Corpus coverage. Of 360 catalogued filings, 282 still download and 263 extract to usable text, leaving 114 of the 150 questions answerable from the local corpus.

2.2 System architecture

The pipeline has four stages: document processing (`src/data_processing`), retrieval (`src/retrieval`), generation (`src/generation`), and evaluation (`src/evaluation`), tied together by `run_experiments.py`. Each stage writes its output to disk so later stages and repeated runs reuse it.

2.3 Document processing

PDFs are downloaded from the catalogue URLs with SEC-compliant headers and retry logic (`download_pdfs.py`). Text is extracted with `pdfplumber`, which recovers both body text and table content, the latter mattering for filings where the answer sits in a financial statement. Extraction is cached per document so that the three

chunk strategies do not each re-parse the same PDF; this turns roughly four and a half hours of repeated extraction into a single ninety-minute pass. Extracted text is cleaned to normalise whitespace and de-hyphenate line breaks before chunking.

2.4 Chunking strategies

We compare three strategies: fixed windows of 256 tokens with 32-token overlap, fixed windows of 512 tokens with 64-token overlap, and a structure-aware strategy that packs whole sentences up to a 512-token limit so a chunk never ends mid-sentence. Overlap in the fixed strategies preserves sentences that straddle a boundary. Each chunk carries its document name, company, sector, period, and a chunk index used to build citation tags.

2.5 Embedding models

The baseline embedding model is BAAI/bge-base-en-v1.5 (768 dimensions). We compare it against all-MiniLM-L6-v2 (384 dimensions), which is smaller and faster, to measure what the larger model buys in retrieval quality. Embeddings are normalised and indexed for inner-product search, which for unit vectors is cosine similarity.

2.6 Retrieval

Chunks are embedded and stored in a FAISS flat inner-product index, one index per embedding model and chunk strategy. At query time the question is embedded with the same model and the top-k chunks are returned. We vary k over 3, 5, and 10. Beyond plain dense search we test three retrieval strategies (Section 4.2): restricting the candidates to the company named in the question, fusing the dense ranking with a BM25 lexical ranking, and reranking. The reranker retrieves a pool of 50 candidates and reorders them with bge-reranker-base, a cross-encoder that scores a question and a passage together rather than embedding them separately, which is slower but sensitive to whether a passage answers the question rather than merely resembling it. It runs locally like the rest of the pipeline.

2.7 Generation

Answers are generated by Qwen3.5-4B, run as 4-bit GGUF weights through llama.cpp with Metal acceleration on Apple silicon. (The brief's suggested Llama-3.1-8B with bitsandbytes 4-bit loading is not usable here, since bitsandbytes has no macOS build; llama.cpp provides the equivalent quantised local inference.) The generator receives the retrieved passages, each prefixed with its citation tag, and is instructed to answer only from them, cite each claim, and say when the passages do not contain the answer. As a second generator we compare against gemma-4-12B, a larger model from a different family. Decoding uses temperature 0.1 for near-deterministic extraction.

2.8 Evaluation metrics

Retrieval is scored with Recall@k, Precision@k, and mean reciprocal rank; a retrieved chunk counts as matching the gold evidence when it covers at least half of the evidence's content words (Section 5.2). Generation is scored with ROUGE-L,

embedding-based semantic similarity, and exact match against the gold answer. Citation quality is precision, recall, and F1 over the [DOCNAME_c<index>] tags, each checked against the passages actually retrieved.

3. Experimental setup

3.1 Implementation

The pipeline is Python 3.12. Retrieval uses sentence-transformers and faiss-cpu; generation uses llama-cpp-python built with Metal; evaluation uses rapidfuzz, rouge-score, and sentence-transformers. Embedding runs on the Apple GPU through MPS. Hardware: Apple M4 Pro, 24 GB unified memory. Full versions are pinned in `requirements.txt`.

3.2 Ablation design

One configuration is the reference (512-token chunks, bge-base embeddings, $k=5$, Qwen3.5-4B), the baseline against which the others are read. Each other run changes one variable so its effect is isolated: chunk strategy (256-token and structure-aware), k (3 and 10), embedding model (MiniLM), generator (gemma-4-12B), retrieval strategy (company metadata filtering, a dense + BM25 hybrid, and cross-encoder reranking over a company-filtered pool), and the generator’s instructions (a prompt that permits deriving figures and showing arithmetic). The experiment matrix is `configs/experiments.yaml`. Indexes are cached and shared across runs, so each embedding-model-and-strategy pair is built once.

Two arms deliberately change two things at once, so a third separates them. Reranking is applied on top of the company filter rather than alone, because the pool the cross-encoder reorders is worth more when it is already restricted to the right company; the filter is therefore also run with generation enabled so its contribution can be measured on its own (Table 2).

4. Results

All figures referenced here are in `results/figures/`, and the full per-configuration tables are in `results/results_summary.md`. Retrieval numbers are on the 114-question answerable subset unless stated; the answerable-versus-all gap is reported separately in Section 4.1.

4.1 Overall performance

The reference configuration (512-token chunks, bge-base embeddings, $k=5$, Qwen3.5-4B) retrieves the gold evidence into its top 5 passages for 45.6% of answerable questions, with a mean reciprocal rank of 0.323. Over all 150 questions, including the 36 whose document is missing or corrupt, $\text{Recall}@5$ falls to 0.360. The gap between 0.360 and 0.456 is almost entirely the missing documents: those questions can never be retrieved, so they enter the full-set average as zeros. This is why we carry both numbers throughout, and Figure 2 shows the gap for every configuration.




Figure 2

Figure 2. Recall@5 over all 150 questions against the 114 answerable. The gap is roughly constant across configurations, which is what we expect if it comes from missing data rather than from any one design choice.

Table 1 gives retrieval quality for every configuration on the answerable subset, ordered by Recall@5.

Table 1. Retrieval metrics on the 114-question answerable subset.

Configuration	Ablation axis	Recall@1	Recall@5	Recall@10	MRR
company filter + reranking	reranking	0.249	0.516	0.516	0.360
company filter	retrieval filter	0.259	0.509	0.509	0.358
structure-aware chunks	chunk size	0.244	0.475	0.475	0.333
512-token chunks (reference)	reference	0.232	0.456	0.456	0.323
k = 10	retrieval k	0.232	0.456	0.504	0.329
derive prompt	prompting	0.232	0.456	0.456	0.323
k = 3	retrieval k	0.232	0.409	0.409	0.311
MiniLM embeddings	embedding model	0.149	0.395	0.395	0.243
dense + BM25 hybrid	retrieval filter	0.124	0.379	0.379	0.218
256-token chunks	chunk size	0.161	0.282	0.282	0.203

The derive-prompt run changes only the generator’s instructions, so its retrieval scores are identical to the reference by construction; it appears here for completeness and is discussed in Section 4.2.

Generation is weaker and limited by retrieval. On the answerable subset the reference system reaches ROUGE-L 0.10 and semantic similarity 0.33, and declines to answer on 102 of 150 questions. The abstention rate follows from retrieval: when the evidence does not reach the model, the prompt tells it to refuse, so a retrieval miss becomes an abstention rather than a wrong figure.

4.2 Ablation analysis

Figure 3 collects the whole ablation into one panel, colouring each configuration by the axis it varies.




Figure 3

Figure 3. Recall@5 by configuration, coloured by ablation axis. The dashed line marks the reference. Company filtering followed by reranking gives the highest score; structure-aware chunking is the best of the content-only changes. The derive-prompt run sits exactly on the reference line because it changes only the generator’s instructions.

Chunk size and structure

Chunking has the largest effect on retrieval of any axis we varied. Cutting the fixed window from 512 to 256 tokens is clearly harmful: Recall@5 drops from 0.456 to 0.282 and MRR from 0.323 to 0.203. Smaller chunks split a table or a paragraph across more pieces, so the evidence for a question is spread thinner and a single retrieved chunk covers less of it. The structure-aware strategy, which packs whole sentences up to the 512-token limit so a chunk never ends mid-sentence, is the best of the three: Recall@5 rises to 0.475 and MRR to 0.333, just above the 512-token baseline, though the margin over it is small.

Embedding model

The larger embedding model helps. Swapping bge-base (768 dimensions) for all-MiniLM-L6-v2 (384 dimensions) lowers Recall@5 from 0.456 to 0.395 and MRR from 0.323 to 0.243, on the same chunks and questions. MiniLM is faster and a third of the dimensionality, but on terse financial questions where the answer is one row of a table, bge-base finds that row more often.

Retrieval k

Increasing k mostly buys recall depth, not rank quality. Moving from $k=5$ to $k=10$ leaves $\text{Recall}@5$ unchanged at 0.456 (the same top-5 passages are retrieved) but lifts $\text{Recall}@10$ to 0.504, so a handful of questions have their evidence sitting in positions 6 through 10. MRR barely moves (0.323 to 0.329), meaning the rank of the first correct passage is stable. At the other end, $k=3$ lowers recall to 0.409 because fewer passages are retrieved. For the generator k is a trade-off: more passages raise the chance the evidence is present but also add distractors to the prompt. Figure 4 shows the recall-versus- k curves.




Figure 4

Figure 4. Recall@k for the reference, the structure-aware chunker, and k=10. Raising k lifts the tail of the curve ($\text{Recall}@10$) without changing where the first correct passage lands.

Retrieval filtering and hybrid search

The content-only changes above move $\text{Recall}@5$ within a narrow band (0.28 to 0.48), which fits the picture that dense similarity over the whole corpus is the ceiling. Two changes attack that ceiling directly. Each FinanceBench question names its target company, and every chunk carries that company in its metadata, so we can restrict the search to the chunks of the company named in the question (inferred from the question text, never the gold answer). This company filter is the single best configuration: $\text{Recall}@5$ rises to 0.509 and MRR to 0.358, above every chunking, k , or embedding change. The gain is bounded by company extraction, a simple name-match catches the company in about 85% of questions, so the remaining misses inherit the unfiltered behaviour; a better entity extractor would raise it further.

Hybrid retrieval, fusing the dense ranking with a BM25 lexical ranking by reciprocal rank fusion, does the opposite: $\text{Recall}@5$ falls to 0.379, below the reference. On these filings BM25 rewards shared boilerplate vocabulary, and the natural-language question terms rarely match the tabular numbers that hold the answer, so the lexical signal adds noise rather than precision. A useful negative result: not every standard retrieval add-on helps on this corpus.

Reranking, and what the retrieval metric cannot see

Recall keeps climbing well past $k=10$: under the company filter it reaches 0.583 at $k=10$, 0.702 at $k=20$ and 0.801 at $k=50$. The evidence is therefore usually inside a deep candidate pool and simply ranked too low, which is what a cross-encoder is for. We retrieve a company-filtered pool of 50 and reorder it with bge-reranker-base, a locally run cross-encoder that scores each question-and-passage pair jointly.

By Recall@5 this looks like almost nothing: 0.509 to 0.516, a gain of 0.007 against an available 0.29. By answer quality it is the largest single improvement in this study. Because gold answers are frequently bare figures, we also score numeric agreement, counting an answer correct when a gold figure appears in it (years excluded, 0.5% tolerance). On that measure reranking takes the system from 0.095 to 0.191, doubling the number of correct answers from 8 to 16, significant by a McNemar exact test ($p = 0.039$). Abstentions fall from 102 to 89 of 150, and the model is also more accurate on the questions it does attempt.

Running the company filter with generation separates the two changes, and the result inverts the retrieval picture:

Table 2. Where the gain actually comes from, answerable subset.

Component	Share of Recall@5 gain	Share of answer-quality gain
company filter	+0.053 (88%)	+0.012 (12%)
cross-encoder reranking	+0.007 (12%)	+0.083 (87%)

The retrieval metric ranks these two interventions in the opposite order to their effect on answers. The reason is visible in a direct measurement: reranking lifts the rate at which the correct filing appears in the top 5 from 0.605 to 0.746, while barely moving the token-overlap score the metric is built on. Getting the right document in front of the generator is what the generator needs, and Recall@5 as defined here is close to blind to it.

Two cases make this concrete. Asked for 3M’s FY2018 capital expenditure (gold \$1,577m), the filtered system answered “\$1,493 million” from a 2015 filing and the reranked system answered “\$1,577 million” from the 2020 filing; both scored Recall@5 of 0.00. Asked for Amazon’s FY2019 net income (gold \$11,588m), both scored Recall@5 of 1.00, but the filtered system answered “\$33,364 million” and the reranked system “\$11,588 million”. Of the nine questions reranking fixed, Recall@5 improved on only three.

Reranking is not uniformly better: it also caused two questions that were previously answered correctly to be refused. The trade is nine gained against two lost.

Prompting and abstention

The system refuses on 102 of 150 questions, and on 40% of questions whose evidence was successfully retrieved, which suggests instructions rather than evidence are the binding constraint for some of them. Inspection supported this: asked for a capital-expenditure figure, the model replied that the excerpts did not contain it “in a cash flow statement format”, a refusal about presentation rather than

absence. We therefore tested a second prompt that states that deriving a figure from inputs present in the excerpts is a valid answer, permits one line of arithmetic, and demotes refusal from the first rule to the last.

It does not work. Refusals move from 102 to 100 of 150 and numeric agreement from 0.095 to 0.119, a difference of two questions ($p = 0.50$). The model largely rephrases its refusals rather than attempting more answers, which also exposed a measurement bug: our abstention counter matched only the exact instructed refusal string, so reworded refusals were being counted as attempts. Correcting it raised the reference system’s abstention count from 99 to 102 and is the reason the figures here differ from an earlier version of this report. Accuracy on the questions the model does answer improves (0.242 to 0.278), so permitting arithmetic plausibly helps once it commits; it simply does not make it commit more often. We report this as a negative result.

Generation model

Making the generator larger does not help. Holding retrieval fixed at the reference setup and swapping Qwen3.5-4B for gemma-4-12B, a model three times the size from a different family, moves none of the generation metrics in gemma’s favour: ROUGE-L is 0.096 against Qwen’s 0.100, semantic similarity is lower (0.253 against 0.329), and citation F1 is level (0.077 against 0.080). gemma only declines more often, on 112 of 150 questions against Qwen’s 102 (Figure 5). Given identical retrieval, the larger model is more conservative about answering from passages that may not contain the fact, which makes little difference to the scores. This fits the rest of the ablation: when the evidence reaches the model less than half the time, a bigger generator has little to work with. gemma-4-12B also runs at roughly a third of Qwen’s tokens per second on the same GPU, so it is the more expensive of the two for no measured benefit.

Table 3. Generation and citation metrics, answerable subset. Numeric agreement counts an answer correct when a gold figure appears in it, and is reported because ROUGE-L cannot score a bare figure: metrics-generated questions score 0.001 on ROUGE-L whatever the answer.

Configuration	Generator	ROUGE-L	Semantic sim	Numeric agr.	Citation F1	Declined (of 150)
reference	Qwen3.5-4B	0.100	0.329	0.095	0.080	102
256-token chunks	Qwen3.5-4B	0.101	0.306	0.107	0.113	105
structure-aware	Qwen3.5-4B	0.102	0.321	0.107	0.070	108
derive prompt	Qwen3.5-4B	0.084	0.329	0.119	0.245	100
company filter	Qwen3.5-4B	0.105	0.355	0.107	0.214	95
company filter + reranking	Qwen3.5-4B	0.116	0.380	0.191	0.130	89
gen_gemma	gemma-4-12B	0.096	0.253	0.095	0.077	112




Figure 5

Figure 5. Qwen3.5-4B against gemma-4-12B on identical retrieval. The larger generator matches or trails on every metric and abstains more.

4.3 Error analysis

Splitting the answerable questions by where they fail is more informative than the aggregate scores. For the reference configuration, the gold evidence reaches the top 5 passages for 56 of 114 answerable questions (49%); the other 58 are retrieval misses, where the model never sees the passage it needs. Among the 56 where retrieval succeeds, the generator still produces a recognisable answer (ROUGE-L above 0.1) only 23 times. So the two failure modes are of similar size, as Figure 6 lays out: roughly half of the answerable questions are lost at retrieval, and of the half that survive retrieval, fewer than half are answered well. Improving either stage alone would leave most of the gap in place.




Figure 6

Figure 6. Where the 114 answerable questions are lost. 51% never retrieve the evidence, 29% retrieve it but answer weakly, and 20% are answered well.

The failures also split by question type (Figure 7). Novel questions are the easiest to retrieve for, with Recall@5 of 0.68, well above the domain-relevant (0.36) and metrics-generated (0.35) types. Metrics-generated questions expose a measurement artifact as much as a model failure: their gold answers are bare figures like “\$1577.00”, so ROUGE-L against the generated sentence is near zero (0.001) even when the number is arguably present, and semantic similarity (0.204) is the more honest signal for them. Domain-relevant questions, whose answers are short sentences rather than single numbers, score highest on semantic similarity (0.455), which fits: lexical and semantic overlap reward sentence answers and penalise bare numbers, regardless of correctness.




Figure 7

Figure 7. Reference-run metrics by question type, answerable subset. Metrics-generated questions score near zero on ROUGE-L because their answers are bare numbers.

Table 4. Reference run by question type (answerable subset).

Question type	Recall@5	MRR	ROUGE-L	Semantic sim	Citation F1
domain-relevant	0.357	0.207	0.154	0.455	0.166
metrics-generated	0.347	0.295	0.001	0.204	0.059
novel-generated	0.681	0.485	0.135	0.307	0.064

5. Discussion

5.1 Key findings

Retrieval is the binding constraint. Among content-only changes the larger embedding model and structure-aware chunking help while halving the chunk beats nothing, but they all stay within a narrow band, which says pure dense search over the whole corpus is close to its ceiling. Two changes raise it, and both use structure the corpus already carries rather than a bigger model: restricting retrieval to the company named in the question, and reordering a deep candidate pool with a cross-encoder. Together they give the best configuration at Recall@5 0.52 and double the number of questions answered correctly. A generic BM25 hybrid instead hurts, because financial filings share boilerplate that the lexical signal over-weights, and a three-times-larger generator on identical retrieval changes nothing, which confirms the constraint is upstream.

The second finding concerns measurement. Recall@5 credits the company filter with 88% of the retrieval gain and reranking with 12%; on whether the system answers correctly those shares reverse. A metric that asks whether one chunk shares half its content words with a gold span cannot see reordering, and reordering is what determines whether the right filing reaches the generator. Tuning against Recall@k

alone would have led us to keep the weaker change and drop the stronger one. ROUGE-L has the same defect on the generation side, scoring 0.001 wherever the gold answer is a bare figure.

Reporting the answerable subset separately matters too: it turns the reference system's Recall@5 from 0.36 over all questions into 0.46 on the questions whose document is present. That gap is data, not model.

5.2 Limitations

Corpus coverage is the largest limitation: a fifth of the filings no longer download and a further set are corrupt, so 36 of the 150 questions ask about a document the system never sees. Reporting the answerable subset keeps that gap apart from model quality, but the smaller subset also makes the per-type numbers noisier.

The retrieval metric is a proxy. It counts a chunk as relevant when it contains at least half of the evidence passage's content words, which handles the length mismatch that breaks exact-span or whole-string matching, but it can still credit a chunk that shares vocabulary without the specific figure, or miss a paraphrase. The 0.5 threshold is defensible rather than tuned.

Generation uses small 4-bit models so the pipeline runs locally and for free, so absolute answer quality trails a larger hosted model, especially on arithmetic over a table. Because the models refuse when the passages do not support an answer, most error surfaces as abstention, the safer failure here but still a ceiling on how many questions get answered. Generation itself is reproducible: re-running a configuration returns bit-identical answers, and so identical ROUGE-L and exact-match scores. An earlier version of this report attributed observed run-to-run drift to the llama.cpp Metal backend, which was wrong. The drift came from our own citation-recall metric, which sampled keywords from a Python set and so depended on per-process hash ordering; the same answer scored differently on each run. That is fixed, but the citation figures reported here predate the fix and are one draw from that distribution, so citation F1 should be read as indicative rather than exact. Citation precision is unaffected and is the reliable citation number. We report a single run per configuration, so small differences between close configurations should still be read with care. Finally, pdfplumber recovers table text but flattens its row and column structure, one likely reason table-heavy numerical questions are hardest.

5.3 Generalisability

The pipeline is not specific to finance. Any closed PDF corpus with question-and-evidence data runs the same way; only the table-aware extraction and citation-tag format are tuned to filings. The methodological point transfers: when a benchmark's source corpus decays, as FinanceBench's links have, an evaluation that does not separate missing data from model error understates the system.

6. Conclusion

We built a RAG question-answering pipeline for financial filings on free, local models, and measured how chunking, k, embedding model, generator, retrieval filtering, reranking and prompting affect performance on FinanceBench. Retrieval sets the ceiling: the gold evidence reaches the model for barely half the answerable

questions, and the content-only knobs move it little. What helped was structure the corpus already carries, not a bigger model. Filtering to the company named in the question and then reordering a deep pool with a cross-encoder together lift Recall@5 from 0.46 to 0.52 and double the number of questions answered correctly, from 8 to 16. A BM25 hybrid hurt, a larger generator did nothing, and a prompt written to make the model refuse less simply made it rephrase its refusals. The small 4-bit generators mostly fail safe, declining rather than inventing figures.

The finding that generalises is about measurement. Ranked by Recall@5 the company filter is worth seven times what reranking is worth; ranked by whether the system answers correctly, reranking is worth seven times the filter. Trusting the retrieval metric would have meant keeping the weaker change and dropping the stronger one.

Three directions follow. An exact-span evidence metric would replace the token-coverage proxy that produced the inversion above. Reranking has more to give, since Recall@50 under the company filter is 0.80 against the 0.52 currently reaching the generator, so a stronger cross-encoder or a deeper pool should convert more of that gap. And recovering the missing corpus would let the full-set numbers measure the model rather than the dataset's decay.

References

Islam, P. et al. (2023). FinanceBench: A New Benchmark for Financial Question Answering. <https://github.com/patronus-ai/financebench>

Appendix

Reproduction instructions are in README.md. Per-configuration metrics and detailed per-question results are in results/experiments/.